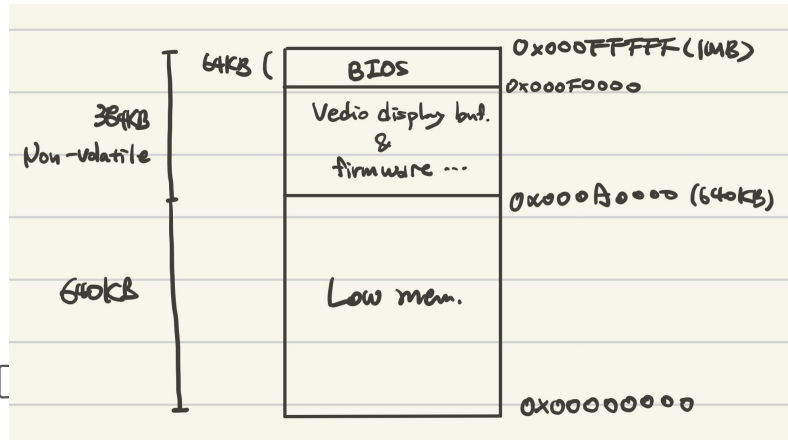


Lab0. Tutorial

1. Physical Address Space

The 1st PC

- 16-bit Intel 8088 processor
- 1MB의 physical memory만 사용할 수 있었다
- 640KB Low memory
 - 당시 PC가 쓸 수 있었던 RAM area (640KB)
 - 16KB, 32KB, 64KB RAM만 가능했다.
- 384KB : reserved by HW (Video display buffers, firmware 등을 위해..)
 - BIOS (Basic Input Output System)
 - jobs: 1. basic system initialization, 2.loading OS 3. pass control of the machine to OS
 - It is stored in



Early PC	Current PC
ROM(Read Only Memory)	updateable flash memory

1. Physical Address Space

After Intel “broke the one megabyte barrier”

hole : RAM 영역을 “Low memory”와 “Extended Memory”로 구분

- PC architects preserves the original layout for the low 1MB of physical address (640KB RAM area)
- in order to ensure backward compatibility with existing software (호환성 유지)

x86 processor가 4GB 이상의 RAM을 지원하기 시작

- 1MB 위로 확장 가능
- 2nd hole이 생기는 이유?

1. Physical Address Space

ROM BIOS


QEMU: virtual machine software

- To boot an OS, we need to hardware environment.
- QEMU allows us to get our own private hardware to boot an OS on.

1. Physical Address Space

ROM BIOS

- IBM PC는 physical addr. **0x000ffff0**에서 **executing**을 시작한다.
 - 0x000ffff0: ROM BIOS를 위해 reserve된 맨 위의 64KB area에 위치.
- PC는 CS = 0xf000, IP = 0xffff에서 **execution**을 시작한다.
- jmp instruction이 가장 처음 실행된다.
 - **jump to CS=0xf000, IP=0xe05b** -> **Q.**
- Physical address = 16 * segment + offset
 - segmented address CS:IP = 0xf000:ffff
 - $16 * 0xf000 + 0xffff = 0xf0000 + 0xffff = \mathbf{0xffff0}$



1. Physical Address Space

Exercise 1.

첫 instruction (jump to 0xf000:205b)이 끝난 후

```
(gdb) si
[f000:e05b] 0xfe05b: cmpw $0xffc8,%cs:(%esi)
0x0000e05b in ?? ()
(gdb) si
[f000:e062] 0xfe062: jne 0xd241d416
0x0000e062 in ?? ()
(gdb) si
[f000:e066] 0xfe066: xor %edx,%edx
0x0000e066 in ?? ()
(gdb) si
[f000:e068] 0xfe068: mov %edx,%ss
0x0000e068 in ?? ()
(gdb) si
[f000:e06a] 0xfe06a: mov $0x7000,%sp
0x0000e06a in ?? ()
(gdb) si
[f000:e070] 0xfe070: mov $0x2d4e,%dx
0x0000e070 in ?? ()
(gdb) si
[f000:e076] 0xfe076: jmp 0x5575ff02
0x0000e076 in ?? ()
```

BIOS가 run할 때,

1. setup interrupt descriptor table
2. initialize the PCI bus and some important devices
3. then search for bootable device (floppy, hard drive, CD-ROM)
4. >> read boot loader and transfers control to it

2. The Boot Loader

sector: disk's minimum transfer unit, 512 bytes

- 읽기, 쓰기 연산이 하나 이상의 **sector** 크기로 이루어진다

boot sector: bootable disk의 첫 번째 sector

- boot loader의 코드가 들어있음

BIOS가 bootable disk를 찾으면,

1. 첫 번째 sector인 boot sector를 메모리에 load한다. (0x7c00 ~ 0x7dff)
2. jmp to CS:IP=0000:7c00
3. control of the machine을 boot loader에게 넘겨준다

2. The Boot Loader

PC의 발전으로 인해 CD-ROM으로부터 boot할 수 있게 됐다.

boot from	Floppy or hard disk	CD-ROM
sector size	512 bytes	2048 bytes

>> 하나의 sector만 가져오는 것이 아니라, 더 큰 boot image를 memory로 load할 수 있게 되었다. (control을 넘겨주기 전에 boot image를 가져올 수 있게 됨)

2. The Boot Loader

Our Boot Loader

- conventional hard drive boot mechanism
- bootasm.S (assembly source file) & bootmain.c (C source file)

2 Main Functions of Boot Loader

1. 프로세서를 real mode 에서 32-bit protected mode로 전환
 - a. real mode: physical addr.를 사용하는 모드. PC가 boot할 때, 몇 개의 instruction을 real mode에서 실행한다.
 - b. 32-bit protected mode
 - i. 이 모드에서만 SW가 (physical) 1MB 위의 메모리에 접근할 수 있다.
 - ii. segmented address를 physical address로 변환하는 과정이 다르다.
 - iii. offset이 16-bit 대신 32-bit이다.
2. x86의 특정 I/O instruction을 통해 IDE disk device register에 직접 접근해서 hard disk로부터 커널을 읽어온다.

Exercise 2

- At what point does the processor start executing 32-bit code? What exactly causes the switch from 16- to 32-bit mode?
- What is the *last* instruction of the boot loader executed, and what is the *first* instruction of the kernel it just loaded?
- *Where* is the first instruction of the kernel?
- How does the boot loader decide how many sectors it must read in order to fetch the entire kernel from disk?
- Where does it find this information?

Exercise 2

Q1. the switch point from 16- to 32-bit mode

```
//PAGEBREAK!  
# Complete the transition to 32-bit protected mode by using a long jmp  
# to reload %cs and %eip. The segment descriptors are set up with no  
# translation, so that the mapping is still the identity mapping.  
ljmp    $(SEG_KCODE<<3), $start32
```

Exercise 2

Q2 ~ Q3 The last instruction of boot loader : call entry()

```
(gdb) x/40i 0x7d3b
0x7d3b: push    %ebp
0x7d3c: mov     %esp,%ebp
0x7d3e: push    %edi
0x7d3f: push    %esi
0x7d40: push    %ebx
0x7d41: sub     $0xc,%esp
0x7d44: push    $0x0
0x7d46: push    $0x1000
0x7d4b: push    $0x10000
0x7d50: call    0x7cf8
0x7d55: add     $0xc,%esp
0x7d58: cmpl    $0x464c457f,0x10000
0x7d62: je      0x7d6c
0x7d64: lea     -0xc(%ebp),%esp
0x7d67: pop     %ebx
0x7d68: pop     %esi
0x7d69: pop     %edi
0x7d6a: pop     %ebp
0x7d6b: ret
```

```
0x7d6c: mov     0x1001c,%eax
0x7d71: lea     0x10000(%eax),%ebx
0x7d77: movzwl 0x1002c,%esi
0x7d7e: shl     $0x5,%esi
0x7d81: add     %ebx,%esi
0x7d83: cmp     %esi,%ebx
0x7d85: jb      0x7d96
0x7d87: call    *0x10018 readseg() in for()
0x7d8d: jmp     0x7d64
0x7d8f: add     $0x20,%ebx
0x7d92: cmp     %ebx,%esi
0x7d94: jbe     0x7d87
0x7d96: mov     0xc(%ebx),%edi
0x7d99: pushl   0x4(%ebx)
0x7d9c: pushl   0x10(%ebx)
0x7d9f: push    %edi
0x7da0: call    0x7cf8 call entry()
0x7da5: mov     0x14(%ebx),%ecx
0x7da8: mov     0x10(%ebx),%eax
0x7dab: add     $0xc,%esp
0x7dae: cmp     %eax,%ecx
```

Exercise 2

Q3 ~ Q4

```
(gdb) x/40i *0x10018
```

```
0x0: push    %ebx
```

1st instruction
of kernel

```
0x1: incl    (%eax)
```

```
0x3: lock push %ebx
```

```
0x5: incl    (%eax)
```

```
0x7: lock ret
```

```
0x9: loop    0xb
```

```
0xb: lock push %ebx
```

```
0xd: incl    (%eax)
```

```
0xf: lock push %ebx
```

```
0x11:        incl    (%eax)
```

```
0x13:        lock push %esp
```

```
0x15:        incl    (%eax)
```

```
0x17:        lock push %ebx
```

```
0x19:        incl    (%eax)
```

```
0x1b:        lock push %ebx
```

```
0x1d:        incl    (%eax)
```

```
0x1f:        lock movsl %ds:(%esi),%es:(%edi)
```

```
0x21:        incb    (%eax)
```

```
0x23:        lock xchg %ebp,%ecx
```

```
0x26:        add     %dh,%al
```

```
0x28:        push    %ds
```

```
0x29:        xlat    %ds:(%ebx)
```

```
// Load each program segment (ignores ph flags).
ph = (struct proghdr*)((uchar*)elf + elf->phoff);
eph = ph + elf->phnum;
for(; ph < eph; ph++){
    pa = (uchar*)ph->paddr;
    readseg(pa, ph->filesz, ph->off);
    if(ph->memsz > ph->filesz)
        stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
}
```

^ Q.

phnum: the number of entries in the program header table.

phoff: pointer of the start of the program header table.

2. The Boot Loader

Loading The Kernel

ELF format : Executable Linkable Format

- C source code들을 컴파일러가 object file로 컴파일하면, 링커가 그 컴파일된 object file들을 하나의 binary image로 만드는데, 이 때 ELF format의 binary로 만든다.
- loading information을 담은 header로도 이해할 수 있다.
- ELF header의 뒤에는 메모리의 특정 주소로 load될 코드나 데이터인 program section이 따라온다.
- Boot loader는 load 되는 데이터나 코드를 수정 없이 메모리에 로드하고 실행을 시작한다.

<https://refspecs.linuxfoundation.org/elf/elf.pdf>

2. The Boot Loader

Loading The Kernel

ELF binary :

- fixed length ELF header
- **variable length** program header
 - load될 program section의 list를 담고 있다.
- program section:
 - .text: 프로그램의 executable instructions
 - .rodata : Read-only data (ex. C 컴파일러에 의해 만들어지는 ASCII string constants)
 - .data : initialized data (ex. 초기화된 global variable: int x = 5)
 - .bss: uninitialized global variable
 - linker가 프로그램의 memory layout을 구성할 때 uninitialized global var.을 위한 공간을 .bss에 reserve한다. 이때 .bss section의 주소와 크기를 기록한다 (contents를 저장할 필요는 X)
 - loader나 프로그램 자체가 해당 global var.들에 0을 할당한다.
 - .data가 바로 뒤에 따라온다. (in memory)

2. The Boot Loader

objdump -h kernel

Load address : 그 섹션이
load되어야 하는 메모리 주소

Link address : 그 섹션이 실행될
메모리 주소

not loaded into memory
by the program loader

Q. LMA of .text
: 00100000

>> Right after the end of BIOS?

tjdp99@fuzzer-master:~/xv6\$ objdump -h kernel

kernel: file format elf32-i386

Sections:		Size	link addr. VMA	load addr. LMA	File off	Algn
Idx	Name					
0	.text	00006ea2	80100000	00100000	00001000	2**4
		CONTENTS, ALLOC, LOAD, READONLY, CODE				
1	.rodata	000009ec	80106ec0	00106ec0	00007ec0	2**5
		CONTENTS, ALLOC, LOAD, READONLY, DATA				
2	.data	00002516	80108000	00108000	00009000	2**12
		CONTENTS, ALLOC, LOAD, DATA				
3	.bss	0000af88	8010a520	0010a520	0000b516	2**5
		ALLOC				
4	.debug_line	000025e8	00000000	00000000	0000b516	2**0
		CONTENTS, READONLY, DEBUGGING				
5	.debug_info	0001051b	00000000	00000000	0000dafa	2**0
		CONTENTS, READONLY, DEBUGGING				
6	.debug_abbrev	00003946	00000000	00000000	0001e019	2**0
		CONTENTS, READONLY, DEBUGGING				
7	.debug_aranges	000003a8	00000000	00000000	00021960	2**3
		CONTENTS, READONLY, DEBUGGING				
8	.debug_str	00000e5e	00000000	00000000	00021d08	2**0
		CONTENTS, READONLY, DEBUGGING				
9	.debug_loc	00005281	00000000	00000000	00022b66	2**0
		CONTENTS, READONLY, DEBUGGING				
10	.debug_ranges	00000700	00000000	00000000	00027de7	2**0
		CONTENTS, READONLY, DEBUGGING				
11	.comment	00000029	00000000	00000000	000284e7	2**0
		CONTENTS, READONLY				

2. The Boot Loader

- Link address와 Load address는 보통 같다.
- The boot loader는 ELF program header 를 사용하여 section을 어떻게 load할 지 결정한다
- program header는 ELF object 의 어느 부분을 어느 주소에 load 할지 특정한다

ph->paddr : 읽어온 program
section들이 저장될 physical
addr.들을 담고 있다.

```
// Load each program segment (ignores ph flags).
ph = (struct proghdr*)((uchar*)elf + elf->phoff);
eph = ph + elf->phnum;
for(; ph < eph; ph++){
    pa = (uchar*)ph->paddr;
    readseg(pa, ph->filesz, ph->off);
    if(ph->memsz > ph->filesz)
        stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
}
```

2. The Boot Loader

objdump -x kernel

```
kernel:      file format elf32-i386
kernel
architecture: i386, flags 0x00000112:
EXEC_P, HAS_SYMS, D_PAGED
start address 0x0010000c
```

Program Header:

```
LOAD off      0x00001000 vaddr 0x80100000 paddr 0x00100000 align 2**12
      filesz 0x000078ac memsz 0x000078ac flags r-x
LOAD off      0x00009000 vaddr 0x80108000 paddr 0x00108000 align 2**12
      filesz 0x00002516 memsz 0x0000d4a8 flags rw-
STACK off     0x00000000 vaddr 0x00000000 paddr 0x00000000 align 2**4
      filesz 0x00000000 memsz 0x00000000 flags rwx
```

memory에
load 되어야
하는 object의
영역

Exercise 3

Boot sector의 load addr.와 link addr. = 0x7c00

0x7c00

```
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) x/40i 0x7c00
=> 0x7c00:    cli
    0x7c01:    xor    %eax,%eax
    0x7c03:    mov    %eax,%ds
    0x7c05:    mov    %eax,%es
    0x7c07:    mov    %eax,%ss
    0x7c09:    in     $0x64,%al
    0x7c0b:    test   $0x2,%al
    0x7c0d:    jne    0x7c09
    0x7c0f:    mov    $0xd1,%al
    0x7c11:    out    %al,$0x64
    0x7c13:    in     $0x64,%al
    0x7c15:    test   $0x2,%al
    0x7c17:    jne    0x7c13
    0x7c19:    mov    $0xdf,%al
    0x7c1b:    out    %al,$0x60
```

0x7c01

Q

```
(gdb) x/40i 0x7c01
=> 0x7c00:    xchg   %ax,%ax
    0x7c02:    nop
    0x7c03:    cli
    0x7c04:    xor    %eax,%eax
    0x7c06:    mov    %eax,%ds
    0x7c08:    mov    %eax,%es
    0x7c0a:    mov    %eax,%ss
    0x7c0c:    in     $0x64,%al
    0x7c0e:    test   $0x2,%al
    0x7c10:    jne    0x7c0c
    0x7c12:    mov    $0xd1,%al
    0x7c14:    out    %al,$0x64
    0x7c16:    in     $0x64,%al
    0x7c18:    test   $0x2,%al
    0x7c1a:    jne    0x7c16
    0x7c1c:    mov    $0xdf,%al
    0x7c1e:    out    %al,$0x60
    0x7c20:    lgdtl  (%esi)
```

Exercise 3

0x7c00

```
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) x/40i 0x7c00
=> 0x7c00:    cli
0x7c01:    xor    %eax,%eax
0x7c03:    mov    %eax,%ds
0x7c05:    mov    %eax,%es
0x7c07:    mov    %eax,%ss
0x7c09:    in     $0x64,%al
0x7c0b:    test   $0x2,%al
0x7c0d:    jne     0x7c09
0x7c0f:    mov    $0xd1,%al
0x7c11:    out     %al,$0x64
0x7c13:    in     $0x64,%al
0x7c15:    test   $0x2,%al
0x7c17:    jne     0x7c13
0x7c19:    mov    $0xdf,%al
0x7c1b:    out     %al,$0x60
```

0x7cf0

Q

```
(gdb) c
Continuing.
```

```
tjdp99@fuzzer-master:~/xv6$ objdump -h bootblock.o
```

```
bootblock.o:      file format elf32-i386
```

```
Sections:
```

Idx	Name	Size	VMA	LMA	File off	Algn
0	.text	000001c0	00007cf0	00007cf0	00000074	2**2

2. The Boot Loader

- Boot loader와 달리, kernel의 load addr.와 link addr.는 다르다
 - kernel은 low addr.에 load되고, high addr.에서 실행된다 (link addr.)
- e_entry : 프로그램의 entry point의 link addr.
 - 프로그램이 실행되어야 하는 text section의 주소 : 0x0010000c **Q. 0x00100000?**

```
tjdp99@fuzzer-master:~/xv6$ objdump -f kernel
```

```
kernel:      file format elf32-i386
architecture: i386, flags 0x00000112:
EXEC_P, HAS_SYMS, D_PAGED
start address 0x0010000c
```

Exercise 4

```
(gdb) b *0x7c00
```

```
Breakpoint 1 at 0x7c00 >> The point the BIOS enters the boot loader
```

```
(gdb) b *0x0010000c
```

```
Breakpoint 2 at 0x10000c >> Boot loader enters the kernel
```

```
(gdb) c
```

```
Continuing.
```

```
[ 0:7c00] => 0x7c00: cli
```

```
Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
```

```
(gdb) x/20x 0x00100000
```

0x100000:	0x00000000	0x00000000	0x00000000	0x00000000
0x100010:	0x00000000	0x00000000	0x00000000	0x00000000
0x100020:	0x00000000	0x00000000	0x00000000	0x00000000
0x100030:	0x00000000	0x00000000	0x00000000	0x00000000
0x100040:	0x00000000	0x00000000	0x00000000	0x00000000

```
(gdb) c
```

```
Continuing.
```

```
The target architecture is assumed to be i386
```

```
=> 0x10000c: mov %cr4,%eax
```

Boot loader가 load한 kernel image

```
tjdp99@fuzzer-master:~/xv
```

```
kernel: file format elf32-i386  
architecture: i386, flags 0x00000000  
EXEC_P, HAS_SYMS, D_PAGED  
start address 0x0010000c
```

```
Thread 1 hit Breakpoint 2, 0x0010000c in ?? ()
```

```
(gdb) x/20x 0x00100000
```

0x100000:	0x1badb002	0x00000000	0xe4524ffe	0x83e0200f
0x100010:	0x220f10c8	0x9000b8e0	0x220f0010	0xc020fd8
0x100020:	0x0100000d	0xc0220f80	0x10b5c0bc	0x2ea0b880
0x100030:	0xe0ff8010	0x90669066	0x90669066	0x90669066
0x100040:	0x53e58955	0x10b5f4bb	0x0cec8380	0x106ec068

3. The Kernel

kernel.ld

```
/* Simple linker script for the JOS kernel.
   See the GNU ld 'info' manual ("info ld") to learn the syntax. */

OUTPUT_FORMAT("elf32-i386", "elf32-i386", "elf32-i386")
OUTPUT_ARCH(i386)
ENTRY(_start)

SECTIONS
{
    /* Link the kernel at this address: "." means the current address */
    /* Must be equal to KERNLINK */
    . = 0x80100000;

    .text : AT(0x100000) {
        *(.text .stub .text.* .gnu.linkonce.t.*)
    }

    PROVIDE(etext = .);    /* Define the 'etext' symbol to this value */
}
```

3. The Kernel

Using virtual memory to work around position dependence

OS kernel을 high virtual addr. (ex. 0xf0100000, 0x80100000)에 link하는 이유는, 유저 프로그램이 사용할 프로세서의 virtual addr.를 남겨두기 위해서이다.

많은 machine들은 addr. 0xf0100000에 대해 physical memory를 가지고 있지 않다.

>> 프로세서의 memory management HW를 사용해, physical addr. 0x00100000을 mapping하는 virtual addr. 0xf0100000을 사용한다.

- 0xf0100000 : link addr. (커널 코드가 실행될 곳)
- 0x00100000 : load addr. (boot loader가 실제로 커널을 올리는 physical mem. (RAM의 1MB 지점 = ROM BIOS의 바로 위)
 - 0x0 이 아니라 0x100000인 이유 : 0xa0000:0x100000이 I/O device를 포함하기 때문

>> 따라서 PC가 적어도 1메가 이상의 physical mem.를 가지고 있어야 한다.

3. The Kernel

- 나중에는 physical memory의 256MB(0x00000000~0x0fffffff)를 mapping할 것이지만, 지금은 우선 4MB를 mapping할 것이다.
 - main.c의 1411~1417 라인에 있는 entrypgdir을 사용

```
1411 __attribute__((__aligned__(PGSIZE)))
1412 pde_t entrypgdir[NPDENTRIES] = {
1413     // Map VA's [0, 4MB) to PA's [0, 4MB)
1414     [0] = (0) | PTE_P | PTE_W | PTE_PS,
1415     // Map VA's [KERNBASE, KERNBASE+4MB) to PA's [0, 4MB)
1416     [KERNBASE >> PDXSHIFT] = (0) | PTE_P | PTE_W | PTE_PS, (flags)
1417 };
1418
```

Handwritten annotations:

- A red arrow points from `NPDENTRIES` to the text "directory entries per directory". Above the arrow is the number "1024".
- Below the flags in line 1416, the following are written in red: "present" under `PTE_P`, "write" under `PTE_W`, "page size" under `PTE_PS`, and "(4096)" below "page size".

3. The Kernel

entry.S가 CR0_PG flag를 설정할 때 까지 memory reference는 linear addr.를 사용한다.
(physical addr.와 같다. - boot.S가 linear addr.에서 physical addr.까지의 identity mapping을 해 주기 때문)

CR0_PG가 설정되면, VM HW에 의해 physical addr.로 변환되는 virtual addr.를 사용한다.

entry_pgdir이 virtual addr.를 physical addr.로 변환해준다

- 0xf0000000 ~ 0xf0400000 >> 0x00000000 ~ 0x00400000
- 0x00000000 ~ 0x00400000 >> 0x00000000 ~ 0x00400000

이 범위에 없는 virtual addr.들은 HW exception을 일으킨다

- interrupt handling이 아직 설정되지 않았기 때문에, QEMU가 machine state를 dump하고 끝낸다.

Exercise 5

```
(gdb) si
=> 0x100025:    mov    %eax,%cr0
0x00100025 in ?? ()
(gdb) x/8x 0x00100000
0x100000:      0x1badb002      0x00000000      0xe4524ffe      0x83e0200f
0x100010:      0x220f10c8      0x9000b8e0      0x220f0010      0xc0200fd8
(gdb) x/8x 0x80100000
0x80100000 <multiboot_header>: 0x00000000      0x00000000      0x00000000      0x00000000
0x80100010 <entry+4>:    0x00000000      0x00000000      0x00000000      0x00000000
(gdb) si
=> 0x100028:    mov    $0x8010b5c0,%esp
0x00100028 in ?? ()
(gdb) x/8x 0x00100000
0x100000:      0x1badb002      0x00000000      0xe4524ffe      0x83e0200f
0x100010:      0x220f10c8      0x9000b8e0      0x220f0010      0xc0200fd8
(gdb) x/8x 0x80100000
0x80100000 <multiboot_header>: 0x1badb002      0x00000000      0xe4524ffe      0x83e0200f
0x80100010 <entry+4>:    0x220f10c8      0x9000b8e0      0x220f0010      0xc0200fd8
```

Exercise 5

After comment out `movl %eax, %cr0`

```
(gdb) si
=> 0x100025:      mov      $0x8010b5c0,%esp
0x00100025 in ?? ()
(gdb) x/8x 0x00100000
0x100000:         0x1badb002      0x00000000      0xe4524ffe      0x83e0200f
0x100010:         0x220f10c8      0x90000b8e0     0x220f0010      0xc0200fd8
(gdb) x/8x 0x80100000
0x80100000 <multiboot_header>:  0x00000000      0x00000000      0x00000000      0x00000000
00
0x80100010 <entry+4>:    0x00000000      0x00000000      0x00000000      0x00000000
(gdb)
```

3. The Kernel

Formatted Printing to the Console

```
#include "types.h"
#include "stat.h"
#include "user.h"
```

```
static void
putc(int fd, char c)
{
    write(fd, &c, 1);
}
```

```
static void
printint(int fd, int xx, int base, int sgn)
{
    static char digits[] = "0123456789ABCDEF";
    char buf[16];
    int i, neg;
    uint x;

    neg = 0;
    if(sgn && xx < 0){
        neg = 1;
        x = -xx;
    } else {
        x = xx;
    }

    i = 0;
    do{
        buf[i++] = digits[x % base];
    }while((x /= base) != 0);
    if(neg)
        buf[i++] = '-';

    while(--i >= 0)
        putc(fd, buf[i]);
}
```

```
// Print to the given fd. Only understands %d, %x, %p, %s.
void
printf(int fd, const char *fmt, ...)
{
    char *s;
    int c, i, state;
    uint *ap;

    state = 0;
    ap = (uint*)(void*)&fmt + 1;
    for(i = 0; fmt[i]; i++){
        c = fmt[i] & 0xff;
        if(state == 0){
            if(c == '%'){
                state = '%';
            } else {
                putc(fd, c);
            }
        } else if(state == '%'){
            if(c == 'd'){
                printint(fd, *ap, 10, 1);
                ap++;
            } else if(c == 'x' || c == 'p'){
                printint(fd, *ap, 16, 0);
                ap++;
            } else if(c == 's'){
                s = (char*)*ap;
                ap++;
                if(s == 0)
                    s = "(null)";
                while(*s != 0){
                    putc(fd, *s);
                    s++;
                }
            } else if(c == 'c'){
                putc(fd, *ap);
                ap++;
            } else if(c == '%'){
                putc(fd, c);
            } else {
                // Unknown % sequence. Print it to draw attention.
                putc(fd, '%');
                putc(fd, c);
            }
            state = 0;
        }
    }
}
```

3. The Kernel

```
7918 static int panicked = 0;
7919
7920 static struct {
7921     struct spinlock lock;
7922     int locking;
7923 } cons;
7924
7925 static void
7926 printint(int xx, int base, int sign)    ex. (300, 10, 1)
7927 {
7928     static char digits[] = "0123456789abcdef";
7929     char buf[16];
7930     int i;
7931     uint x;
7932
7933     if(sign && (sign = xx < 0))
7934         x = -xx;
7935     else
7936         x = xx;
7937
7938     i = 0;
7939     do{
7940         buf[i++] = digits[x % base];
7941     }while((x /= base) != 0);
7942
7943     if(sign)
7944         buf[i++] = '-';
7945
7946     while(--i >= 0)
7947         consputc(buf[i]);
7948 }
7949
```

buf 0 0 3 -

- 3 0 0

```
7950 // Print to the console. only understands %d, %x, %p, %s.
7951 void
7952 cprintf(char *fmt, ...)
7953 {
7954     int i, c, locking;
7955     uint *argp;
7956     char *s;
7957
7958     locking = cons.locking;
7959     if(locking)
7960         acquire(&cons.lock);
7961
7962     if (fmt == 0)
7963         panic("null fmt");
7964
7965     argp = (uint*)(void*)&fmt + 1;
7966     for(i = 0; (c = fmt[i] & 0xff) != 0; i++){
7967         if(c != '%'){
7968             consputc(c);
7969             continue;
7970         }
7971         c = fmt[++i] & 0xff; → % 다음 숫자.
7972         if(c == 0) (NULL char.)
7973             break;
7974         switch(c){
7975             case 'd':
7976                 printint(*argp++, 10, 1); → 003을 10진법 print
7977                 break;
7978             case 'x':
7979             case 'p':
7980                 printint(*argp++, 16, 0); → 16진법
7981                 break;
7982             case 's':
7983                 if((s = (char*)*argp++) == 0)
7984                     s = "(null)";
7985                 for(; *s; s++)
7986                     consputc(*s);
7987                 break;
7988             case '%':
7989                 consputc('%'); → % 처리.
7990                 break;
7991             default:
7992                 // Print unknown % sequence to draw attention.
7993                 consputc('%');
7994                 consputc(c);
7995                 break;
7996         }
7997     }
7998 }
```

error.

```

8004 void
8005 panic(char *s)
8006 {
8007     int i;
8008     uint pcs[10];
8009
8010     cli();
8011     cons.locking = 0;
8012     cprintf("cpu with apicid %d: panic: ", cpu->apicid);
8013     cprintf(s);
8014     cprintf("\n");
8015     getcallerpcs(&s, pcs);
8016     for(i=0; i<10; i++)
8017         cprintf(" %p", pcs[i]);
8018     panicked = 1; // freeze other CPU
8019     for(;;)
8020         ;
8021 }

```

2201

apic: advanced
Interrupt
controller


```

8050 #define BACKSPACE 0x100
8051 #define CRTPORT 0x3d4
8052 static ushort *crt = (ushort*)P2V(0xb8000); // CGA memory

```

→ graphic dev.

→ 0xb8000 + KERNBASE
(physical to virtual)



```

8053
8054 static void
8055 cgaputc(int c)
8056 {
8057     int pos;
8058
8059     // Cursor position: col + 80*row.
8060     outb(CRTPORT, 14);
8061     pos = inb(CRTPORT+1) << 8;
8062     outb(CRTPORT, 15);
8063     pos |= inb(CRTPORT+1);

```

pos = pos | inb(CRTPORT+1);

```

8064
8065     if(c == '\n')
8066         pos += 80 - pos%80;
8067     else if(c == BACKSPACE){
8068         if(pos > 0) --pos;
8069     } else

```

crt[pos++] = (c&0xff) | 0x0700; // black on white

Q. (print)

```

8070
8071
8072     if(pos < 0 || pos > 25*80)
8073         panic("pos under/overflow");
8074
8075     if((pos/80) >= 24){ // Scroll up.
8076         memmove(crt, crt+80, sizeof(crt[0])*23*80);
8077         pos -= 80;
8078         memset(crt+pos, 0, sizeof(crt[0])*(24*80 - pos));
8079     }

```

→ overflow →

```


8080
8081     outb(CRTPORT, 14);
8082     outb(CRTPORT+1, pos>>8);
8083     outb(CRTPORT, 15);
8084     outb(CRTPORT+1, pos);
8085     crt[pos] = ' ' | 0x0700;

```

Q.

```
8100 void
8101 consputc(int c)
8102 {
8103     if(panicked){
8104         cli();
8105         for(;;)
8106             ;
8107     }
8108
8109     if(c == BACKSPACE){
8110         uartputc('\b'); uartputc(' '); uartputc('\b');
8111     } else
8112         uartputc(c);
8113     cgaputc(c);
8114 }
```


Calling convention



첫 번째 argument	<pre>state = 0; ap = (uint*)(void*)&fmt + 1;</pre>	<pre>ap = (uint*)(void*)&fmt - 1;</pre>
다음 argument	<pre>ap++;</pre>	<pre>ap -- ;</pre>

Exercise 6

Determine where the kernel initializes its stack, and exactly where in memory its stack is located. How does the kernel reserve space for its stack? And at which “end” of this reserved area is the stack pointer initialized to point to?

After jump to entry point from `entry()` in `bootmain.c`,

Through setting `cr0` which is control register, `CR0_PG` means paging is enabled

```
#define CR0_WP 0x00010000
#define CR0_PG 0x80000000
```

```
#define NPROC 64 // maximum number
#define KSTACKSIZE 4096 // size of per-pro
```

entry:

Turn on page size extension for 4Mbyte pages

```
movl %cr4, %eax
```

```
orl $(CR4_PSE), %eax
```

```
movl %eax, %cr4
```

Set page directory

```
movl $(V2P_W0(entrypgdir)), %eax
```

```
movl %eax, %cr3
```

Turn on paging.

```
movl %cr0, %eax
```

```
orl $(CR0_PG|CR0_WP), %eax
```

```
movl %eax, %cr0
```

Set up the stack pointer.

```
movl $(stack + KSTACKSIZE), %esp
```

Jump to main(), and switch to executing at

high addresses. The indirect call is needed because

the assembler produces a PC-relative instruction

for a direct jump.

```
mov $main, %eax
```

```
jmp *%eax
```

```
.comm stack, KSTACKSIZE
```

(gdb) x/20i 0x0010000c

```
=> 0x10000c: mov %cr4,%eax
0x10000f: or $0x10,%eax
0x100012: mov %eax,%cr4
0x100015: mov $0x109000,%eax
0x10001a: mov %eax,%cr3
0x10001d: mov %cr0,%eax
0x100020: or $0x80010000,%eax
0x100025: mov %eax,%cr0
0x100028: mov $0x8010b5c0,%esp
0x10002d: mov $0x80102ea0,%eax
0x100032: jmp *%eax
0x100034: xchg %ax,%ax
0x100036: xchg %ax,%ax
0x100038: xchg %ax,%ax
0x10003a: xchg %ax,%ax
0x10003c: xchg %ax,%ax
0x10003e: xchg %ax,%ax
0x100040: push %ebp
0x100041: mov %esp,%ebp
0x100043: push %ebx
```

- CR0_PG를 CR0에 넣으면 Paging HW를 enable한다
 - 그러나 여전히 low addr.에서 실행하고 있는 상태이다.
- 커널을 High mem.에 mapping하기 위해
 - esp (stack pointer)를 initialize한다 (movl \$(stack + KSTACKSIZE), %esp)
 - high addr.의 커널 스택이 valid
- main으로 jump한다.
 - Q. p.23 “indirect jump”가 필요하다

+) Stack

- esp : stack pointer, 현재 사용되는 stack의 가장 낮은 지점을 point
- ebp : base pointer
 - SW convention을 따르는 register.
 - C 함수의 entry에서, 함수의 prologue code가 직전 함수의 base pointer(ebp)를 스택에 저장(push)
 - 현재의 esp를 ebp에 copy
 - >> stack에 저장된 ebp를 사용해 nested call을 backtrace할 수 있다!

Assembly trap handlers (p.42)

int instruction : switch the processor from user mode to kernel mode by generating trap

xv6 는 x86 hardware가 int instruction을 만났을 때 handling할 수 있게 set up 한다

tvinit(3317): x86 hardware가 256 interrupts를 가지고 있어서, 256개의 entry를 가진 IDT(interrupt descriptor table)를 set up한다